

Data Exploration and Analysis - Task 1

Siddharth Gada

```
## ## Chunk: Loading Libraries
```

```
library(data.table)
library(ggplot2)
library(ggmosaic)
library(readr)
library(readxl)
library(stringr)
library(dplyr)
library(lubridate)
library(scales)
options(scipen = 999) # Turn off scientific notation
```

```
## ## Chunk: Loading Datasets
```

```
transactionData <- read_excel(paste0(
  "C:/Users/gadas/OneDrive/Desktop/",
  "Classes Outside UNT/Forage Project/",
  "Quantium Data Analytics/QVI_transaction_data.xlsx"))

customerData <- read_csv(paste0(
  "C:/Users/gadas/OneDrive/Desktop/",
  "Classes Outside UNT/Forage Project/",
  "Quantium Data Analytics/QVI_purchase_behaviour.csv"))

head(transactionData)
```

```
## # A tibble: 6 x 8
##   DATE STORE_NBR LYLTY_CAR~1 TXN_ID PROD_~2 PROD_~3 PROD_~4
##   <dbl>   <dbl>   <dbl> <dbl>   <dbl> <chr>   <dbl>
## 1 43390     1     1000     1       5 Natura~     2
## 2 43599     1     1307    348     66 CCs Na~     3
## 3 43605     1     1343    383     61 Smiths~     2
## 4 43329     2     2373    974     69 Smiths~     5
## 5 43330     2     2426   1038    108 Kettle~     3
## 6 43604     4     4074   2982     57 Old El~     1
## # ... with 1 more variable: TOT_SALES <dbl>, and
## #   abbreviated variable names 1: LYLTY_CARD_NBR,
## #   2: PROD_NBR, 3: PROD_NAME, 4: PROD_QTY
```

```
head(customerData)
```

```
##      LYLTY_CARD_NBR      LIFESTAGE PREMIUM_CUSTOMER
## 1          1000  YOUNG SINGLES/COUPLES      Premium
## 2          1002  YOUNG SINGLES/COUPLES      Mainstream
## 3          1003      YOUNG FAMILIES      Budget
## 4          1004  OLDER SINGLES/COUPLES      Mainstream
## 5          1005  MIDAGE SINGLES/COUPLES      Mainstream
## 6          1007  YOUNG SINGLES/COUPLES      Budget
```

```
## ## Data Exploration Start
```

```
## ## Examining Transaction Data
```

```
# Examine date variable from transaction data
summary(transactionData$DATE)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  43282  43373   43464   43464   43555   43646
```

```
# Need to convert date variable from numeric to date format
transactionData$DATE_Converted <- as.Date(transactionData$DATE,
                                           origin = "1899-12-30")
```

```
# Remove DATE column which is not converted
transactionData$DATE <- NULL
```

```
# Examine PROD_NAME
summary(transactionData$PROD_NAME)
```

```
##      Length      Class      Mode
##  264836 character character
```

```
# Split all product names into words
words <- unlist(strsplit(transactionData$PROD_NAME, "\\s+"))
```

```
# Keep only words with letters a-z or A-Z (remove digits/special chars)
clean_words <- words[!grepl("[^A-Za-z]", words)]
```

```
# Create frequency table
word_freq <- data.table(word = clean_words)[, .N, by = word]
setorder(word_freq, -N)
```

```
# Remove salsa products because we are only interested in keeping the
# data related to chips sales
transactionData <- subset(transactionData, !grepl("salsa", tolower(PROD_NAME)))
```

```
# Summarise the data to check for nulls and possible outliers
summary(transactionData)
```

```
##      STORE_NBR      LYLTY_CARD_NBR      TXN_ID
##  Min.   : 1.0   Min.   : 1000   Min.   : 1
## 1st Qu.: 70.0   1st Qu.: 70015   1st Qu.: 67569
```

```
## Median :130.0    Median : 130367    Median : 135183
## Mean   :135.1    Mean   : 135531    Mean   : 135131
## 3rd Qu.:203.0    3rd Qu.: 203084    3rd Qu.: 202654
## Max.   :272.0    Max.   :2373711    Max.   :2415841
##      PROD_NBR      PROD_NAME      PROD_QTY
## Min.    : 1.00    Length:246742    Min.    : 1.000
## 1st Qu.: 26.00    Class :character 1st Qu.: 2.000
## Median : 53.00    Mode  :character Median : 2.000
## Mean    : 56.35                      Mean    : 1.908
## 3rd Qu.: 87.00                      3rd Qu.: 2.000
## Max.    :114.00                      Max.    :200.000
##      TOT_SALES      DATE_Converted
## Min.    : 1.700    Min.    :2018-07-01
## 1st Qu.: 5.800    1st Qu.:2018-09-30
## Median : 7.400    Median :2018-12-30
## Mean    : 7.321    Mean    :2018-12-30
## 3rd Qu.: 8.800    3rd Qu.:2019-03-31
## Max.    :650.000    Max.    :2019-06-30
```

```
data.frame(
  Column = names(transactionData),
  Total_Obs = nrow(transactionData),
  Non_Missing = colSums(!is.na(transactionData)),
  Missing = colSums(is.na(transactionData))
)
```

```
##              Column Total_Obs Non_Missing Missing
## STORE_NBR      STORE_NBR    246742      246742      0
## LYLTY_CARD_NBR LYLTY_CARD_NBR 246742      246742      0
## TXN_ID          TXN_ID      246742      246742      0
## PROD_NBR        PROD_NBR    246742      246742      0
## PROD_NAME       PROD_NAME    246742      246742      0
## PROD_QTY        PROD_QTY    246742      246742      0
## TOT_SALES       TOT_SALES    246742      246742      0
## DATE_Converted DATE_Converted 246742      246742      0
```

```
# There are no nulls in the columns but product quantity appears to have an outlier
# Filter the dataset to find the outlier in product quantity
# where 200 chip packets were bought in one transaction
setDT(transactionData)
transactionData[PROD_QTY==200]
```

```
##      STORE_NBR LYLTY_CARD_NBR TXN_ID PROD_NBR
## 1:          226          226000 226201      4
## 2:          226          226000 226210      4
##              PROD_NAME PROD_QTY TOT_SALES
## 1: Dorito Corn Chp    Supreme 380g    200    650
## 2: Dorito Corn Chp    Supreme 380g    200    650
##      DATE_Converted
## 1:      2018-08-19
## 2:      2019-05-20
```

```
# There are two transactions where 200 packets of chips are bought in  
# one transaction and both of these transactions were by the same customer.
```

```
# Checking if the customer has had other transactions  
transactionData[LYLTY_CARD_NBR==226000]
```

```
##      STORE_NBR LYLTY_CARD_NBR TXN_ID PROD_NBR  
## 1:      226      226000 226201      4  
## 2:      226      226000 226210      4  
##      PROD_NAME PROD_QTY TOT_SALES  
## 1: Dorito Corn Chp Supreme 380g      200      650  
## 2: Dorito Corn Chp Supreme 380g      200      650  
##      DATE_Converted  
## 1:      2018-08-19  
## 2:      2019-05-20
```

```
# Looks like this customer has only had the two transactions over the year  
# and is not an ordinary retail customer.  
# The customer might be buying chips for commercial purposes instead.  
# We'll remove this loyalty card number from further analysis.  
transactionData <- transactionData[LYLTY_CARD_NBR != 226000]
```

```
# Checking if all transactions from the card number have been removed  
transactionData[LYLTY_CARD_NBR==226000]
```

```
## Empty data.table (0 rows and 8 cols): STORE_NBR,LYLTY_CARD_NBR,TXN_ID,PROD_NBR,PROD_NAME,PROD_QTY...
```

```
# Count the number of transactions by date  
transactions_by_date <- transactionData %>%  
  group_by(DATE_Converted) %>%  
  summarise(Transaction_Count = n())
```

```
# There's only 364 rows, meaning only 364 dates which indicates a missing date.  
# Create a full sequence of dates from 1 Jul 2018 to 30 Jun 2019 and use this to  
# create a chart of number of transactions over time to find the missing date.  
ALLDATES <- data.table(DATE_Converted = seq(as.Date("2018-07-01"),  
                                           as.Date("2019-06-30"), by = "day"))
```


```
# Join all_dates with transactions_by_date (left join)  
fullTransactionData <- merge(ALLDATES, transactionData, by = "DATE_Converted",  
                             all.x = TRUE)
```

```
# Count the number of transactions by date to see if the missing date was added  
transactions_by_date_1 <- fullTransactionData %>%  
  group_by(DATE_Converted) %>%  
  summarise(Transaction_Count = n())
```

```
# Setting plot themes to format graphs  
theme_set(theme_bw())  
theme_update(plot.title = element_text(hjust = 0.5))
```

```
# Plot transactions over time
```

```
ggplot(transactions_by_date_1, aes(x = DATE_Converted, y = Transaction_Count))+
  geom_line() +
  labs(x = "Day", y = "Number of Transactions", title = "Transactions over time")+
  scale_x_date(breaks = "1 month")+
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5))
```




```
# We can see that there is an increase in purchases in December and a break
# in late December. Let's zoom in on this.

# Weekly transaction plot for Presentation
fullTransactionData$WeekStart <- floor_date(fullTransactionData$DATE_Converted,
                                             unit = "week", week_start = 1)

weekly_transactions <- fullTransactionData[, .(
  weekly_transactions = uniqueN(TXN_ID) # Count of unique transactions
), by = WeekStart][order(WeekStart)]

ggplot(weekly_transactions, aes(x = WeekStart, y = weekly_transactions))+
  geom_line(color = "steelblue", size = 1.5) +
  labs(x = "Month", y = "Number of Transactions",
       title = "Weekly Transactions") +
  scale_x_date(date_breaks = "1 month", date_labels = "%b",
              limits = as.Date(c("2018-07-01", "2019-06-30")))+
  theme(axis.text.x = element_text(vjust = 0.5))
```




```
fullTransactionData$WeekStart <- NULL
weekly_transactions <- NULL

# Filter to December and look at individual days
# Filter for December 2018
setDT(transactions_by_date_1)
december_data <- transactions_by_date_1[
  transactions_by_date_1$DATE_Converted >= as.Date("2018-12-01")
  & transactions_by_date_1$DATE_Converted <= as.Date("2018-12-31")
]

# Plot daily transactions for December
ggplot(december_data, aes(x = DATE_Converted, y = Transaction_Count))+
  geom_line(color = "darkred")+
  labs(
    x = "Date",
    y = "Number of Transactions",
    title = "Daily Transactions in December 2018"
  ) +
  scale_x_date(
    date_breaks = "1 day",
    date_labels = "%d-%b"
  ) +
  theme(
    axis.text.x = element_text(angle = 90, vjust = 0.5),
    panel.background = element_blank(),
    panel.grid.major = element_line(color = "grey90")
  )
```

Daily Transactions in December 2018



*# We can see that the increase in sales occurs in the lead-up to Christmas and
that there are zero sales on Christmas day itself which was the missing date.
This is due to shops being closed on Christmas day.*

*# Creating a Pack size variable
We can work this out by taking the digits that are in PROD_NAME*
`setDT(transactionData)`
`transactionData[, PACK_SIZE := parse_number(PROD_NAME)]`


Checking if the pack sizes look sensible
`transactionData[, .N, PACK_SIZE][order(PACK_SIZE)]`

##	PACK_SIZE	N
## 1:	70	1507
## 2:	90	3008
## 3:	110	22387
## 4:	125	1454
## 5:	134	25102
## 6:	135	3257
## 7:	150	40203
## 8:	160	2970
## 9:	165	15297
## 10:	170	19983
## 11:	175	66390
## 12:	180	1468
## 13:	190	2995
## 14:	200	4473
## 15:	210	6272

```
## 16:      220  1564
## 17:      250  3169
## 18:      270  6285
## 19:      330 12540
## 20:      380  6416
```

The largest size is 380g and the smallest size is 70g - seems sensible

```
# Histogram of PACK_SIZE
# Treat PACK_SIZE as a factor (categorical)
ggplot(transactionData, aes(x = factor(PACK_SIZE)))+
  geom_bar(fill = "steelblue")+
  labs(
    x = "Pack Size",
    y = "Count",
    title = "Distribution of Pack Size"
  )+
  theme_minimal()
```



```
# Creating a Brand variable
# We can work this out by taking the first word that is in PROD_NAME
transactionData[, Brand := ifelse(is.na(PROD_NAME), NA,
                                  tstrsplit(PROD_NAME, " ")[[1]])]

# Finding out unique Brand names
unique(transactionData$Brand)
```

```
## [1] "Natural"    "CCs"        "Smiths"     "Kettle"
## [5] "Grain"      "Doritos"    "Twisties"   "WW"
```



```
## [9] "Thins"      "Burger"      "NCC"         "Cheezels"
## [13] "Infzns"     "Red"         "Pringles"    "Dorito"
## [17] "Infuzions"  "Smith"       "GrnWves"     "Tyrrells"
## [21] "Cobs"       "French"      "RRD"         "Tostitos"
## [25] "Cheetos"    "Woolworths" "Snbts"       "Sunbites"
```

```
# Clean Brand names
```

```
transactionData[Brand == "Red", Brand := "RRD"]
transactionData[Brand == "Smith", Brand := "Smiths"]
transactionData[Brand == "Infzns", Brand := "Infuzions"]
transactionData[Brand == "Snbts", Brand := "Sunbites"]
transactionData[Brand == "WW", Brand := "Woolworths"]
transactionData[Brand == "NCC", Brand := "Natural"]
transactionData[Brand == "Dorito", Brand := "Doritos"]
transactionData[Brand == "Grain", Brand := "GrnWves"]
```

```
# Checking if any more discrepancies in names
```

```
unique(transactionData$Brand)
```

```
## [1] "Natural"    "CCs"        "Smiths"     "Kettle"
## [5] "GrnWves"    "Doritos"    "Twisties"   "Woolworths"
## [9] "Thins"      "Burger"     "Cheezels"   "Infuzions"
## [13] "RRD"        "Pringles"   "Tyrrells"   "Cobs"
## [17] "French"     "Tostitos"   "Cheetos"    "Sunbites"
```

```
# Clean each PROD_NAME by removing words that contain non-letter characters
```

```
transactionData[, PROD_NAME := sapply(strsplit(PROD_NAME, "\\s+"),
                                         function(words) {
                                           clean_words <- words[!grepl("[^A-Za-z]", words)]
                                           paste(clean_words, collapse = " ")
                                         }
)]
```

```
## ## Examining Customer Data
```

```
# Making sure there are no missing values
```

```
summary(customerData)
```

```
## LYLTY_CARD_NBR      LIFESTAGE      PREMIUM_CUSTOMER
## Min.   : 1000      Length:72637      Length:72637
## 1st Qu.: 66202     Class :character   Class :character
## Median : 134040     Mode  :character   Mode  :character
## Mean   : 136186
## 3rd Qu.: 203375
## Max.   : 2373711
```

```
length(customerData$LYLTY_CARD_NBR)
```

```
## [1] 72637
```

```
length(unique(transactionData$LYLTY_CARD_NBR))
```

```
## [1] 71287
```

```
# Performing a left join because we only want matches from 'transactionData'  
# since 'customerData' has customer 226000 which we excluded from our dataset  
# because he buys for commercial use  
# This is the reason we did not want to delete the customers that bought SALSA  
# so that we get exact matches from the two datasets.
```

```
data <- merge(transactionData, customerData, all.x = TRUE)
```

```
length(unique(data$LYLTY_CARD_NBR))
```

```
## [1] 71287
```

```
# Check if some customers were not matched on by checking for nulls.  
sum(is.na(data$LYLTY_CARD_NBR))
```

```
## [1] 0
```

```
sum(is.na(data$LIFESTAGE))
```

```
## [1] 0
```

```
sum(is.na(data$PREMIUM_CUSTOMER))
```

```
## [1] 0
```

```
# Code to save dataset as a csv  
write.csv(data,  
          file = paste0(  
            "C:/Users/gadas/OneDrive/Desktop/",  
            "Classes Outside UNT/Forage Project/",  
            "Quantium Data Analytics/QVI_data.csv"))
```

```
## ## Data analysis on customer segments Start
```

```
## ## Chunk: Total sales by Loyalty Card Number
```

```
data[, .(Total_Spend = sum(TOT_SALES)), by = .(LYLTY_CARD_NBR)][order(-Total_Spend)]
```

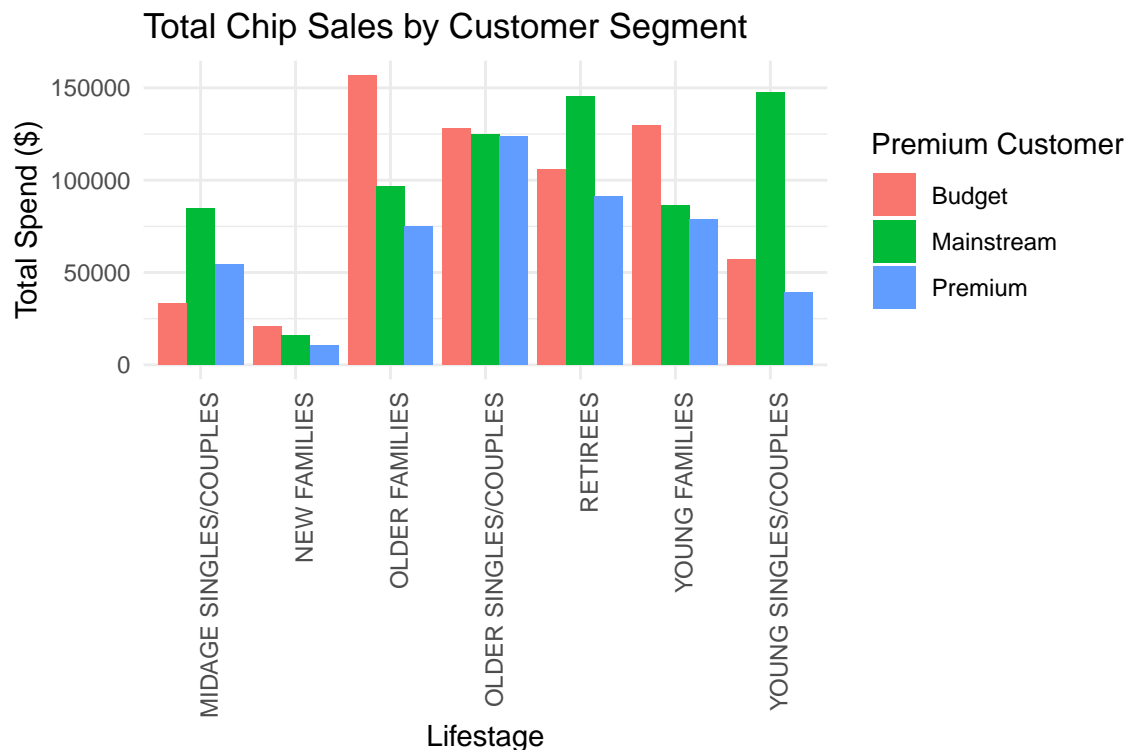
```
##      LYLTY_CARD_NBR Total_Spend  
## 1:      230078      138.6  
## 2:      58361      124.8  
## 3:      63197      122.6  
## 4:     162039      121.6  
## 5:     179228      120.8
```

```
## ---
## 71283:      268247      1.7
## 71284:      268313      1.7
## 71285:      268315      1.7
## 71286:      268390      1.7
## 71287:      268476      1.7

## ## Chunk: Total sales by Lifestage and Premium Customer

# LIFESTAGE: Customer attribute that identifies whether a customer has a family
# or not and what point in life they are at
# PREMIUM_CUSTOMER: Customer segmentation used to differentiate shoppers by the
# price point of products they buy and the types of products they buy.
sales_by_segment <- data[, .(Total_Spend = sum(TOT_SALES)),
                           by = .(PREMIUM_CUSTOMER, LIFESTAGE)][order(-Total_Spend)]

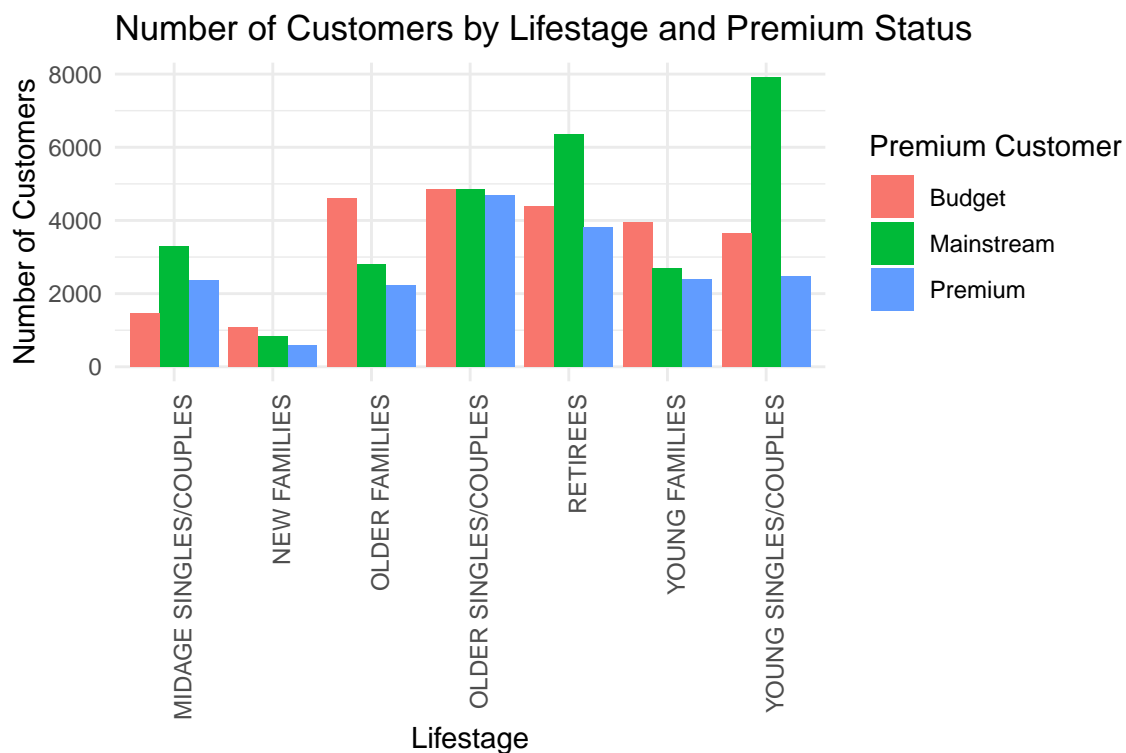
# Plot the sales split
ggplot(sales_by_segment, aes(x = LIFESTAGE, y = Total_Spend, fill = PREMIUM_CUSTOMER))+
  geom_bar(stat = "identity", position = "dodge")+
  labs(title = "Total Chip Sales by Customer Segment",
       x = "Lifestage", y = "Total Spend ($)",
       fill = "Premium Customer")+
  theme_minimal()+
  theme(axis.text.x = element_text(angle = 90, hjust = 1))
```



```
# Sales are coming mainly from Budget - older families,
# Mainstream - young singles/couples and Mainstream - retirees
```

```
## ## Chunk: How many customers are in each segment
```

```
# How many customers are in each segment
customers_by_segment <- data[, uniqueN(LYLT_CARD_NBR),
                                by = .(LIFESTAGE, PREMIUM_CUSTOMER)]
setnames(customers_by_segment, "V1", "N")
# Plot number of customers
ggplot(customers_by_segment, aes(x = LIFESTAGE, y = N, fill = PREMIUM_CUSTOMER))+
  geom_bar(stat = "identity", position = "dodge")+
  labs(title = "Number of Customers by Lifestage and Premium Status",
       x = "Lifestage", y = "Number of Customers",
       fill = "Premium Customer")+
  theme_minimal()+
  theme(axis.text.x = element_text(angle = 90, hjust = 1))
```



```
# There are more Mainstream - young singles/couples and Mainstream - retirees who buy chips.
# This contributes to there being more sales to these customer segments but this
# is not a major driver for the Budget - Older families segment.

# Percentage of Customers by Lifestage and Premium Status plot for Presentation
customers_by_segment_pct <- customers_by_segment %>%
  group_by(LIFESTAGE) %>%
  mutate(pct = N / sum(N),
         pct = pct / sum(pct),
         total_N = sum(N)) %>%
  ungroup()

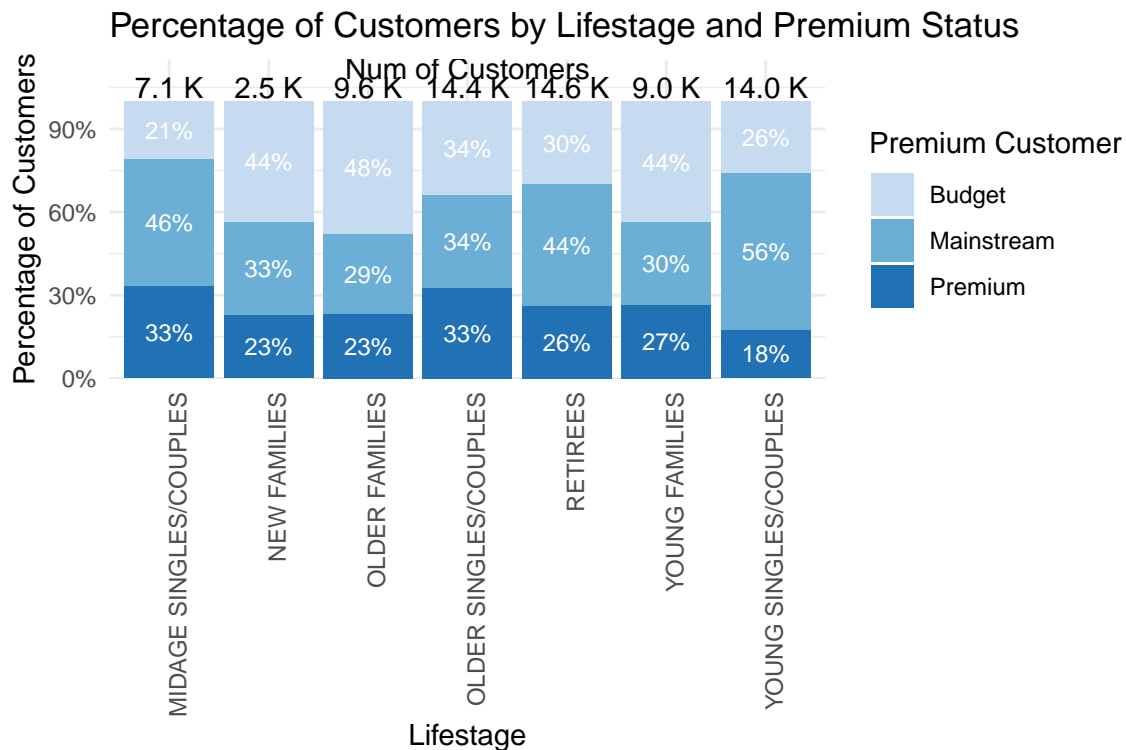
ggplot(customers_by_segment_pct, aes(x = LIFESTAGE, y = pct, fill = PREMIUM_CUSTOMER))+
  geom_bar(stat = "identity", position = "stack") +
  geom_text(aes(label = scales::percent(pct, accuracy = 1)),
```

```

    position = position_stack(vjust = 0.5),
    size = 3, color = "white") +
geom_text(data = customers_by_segment_pct %>%
  distinct(LIFESTAGE, total_N),
  aes(x = LIFESTAGE, y = 1.05,
      label = paste(scales::number(total_N / 1000, accuracy = 0.1), "K")),
  inherit.aes = FALSE) +
annotate("text", x = length(unique(customers_by_segment_pct$LIFESTAGE)) / 2 + 0.5,
  y = 1.12, label = "Num of Customers") +
labs(title = "Percentage of Customers by Lifestage and Premium Status",
  x = "Lifestage", y = "Percentage of Customers",
  fill = "Premium Customer")+
scale_y_continuous(labels = percent_format(), limits = c(0, 1.15), expand = c(0,0)) +
scale_fill_manual(values = c("Budget" = "#c6dbef", # light blue
  "Mainstream" = "#6baed6", # medium blue
  "Premium" = "#2171b5")) + # dark blue

theme_minimal()+
theme(axis.text.x = element_text(angle = 90, hjust = 1))

```



```
customers_by_segment_pct <- NULL
```

```
## ## Chunk: How many chips are bought per customer by segment
```

```


avg_units_per_customer <- data[, .(Num_Chips_bought = sum(PROD_QTY),
  Unique_Customers = uniqueN(LYLT_CARD_NBR)),
  by = .(PREMIUM_CUSTOMER, LIFESTAGE)][order(-Num_Chips_bought)]
# Compute average units per customer

```

```

avg_units_per_customer[, Avg_Units := Num_Chips_bought / Unique_Customers]
# Plot
ggplot(avg_units_per_customer, aes(x = LIFESTAGE, y = Avg_Units, fill = PREMIUM_CUSTOMER))+
  geom_bar(stat = "identity", position = "dodge")+
  labs(title = "Average Chip Units per Customer by Segment",
       x = "Lifestage", y = "Average Units",
       fill = "Premium Customer")+
  theme_minimal()+
  theme(axis.text.x = element_text(angle = 90, hjust = 1))

```




```

# Older families and young families in general buy more chips per customer

# Plot for Presentation
ggplot(avg_units_per_customer, aes(x = LIFESTAGE, y = Avg_Units, fill = PREMIUM_CUSTOMER))+
  geom_bar(stat = "identity", position = "dodge")+
  labs(title = "Average Chip Units per Customer by Segment",
       x = "Lifestage", y = "Average Units",
       fill = "Premium Customer")+
  scale_fill_manual(values = c("Budget" = "#c6dbef", # light blue
                               "Mainstream" = "#6baed6", # medium blue
                               "Premium" = "#2171b5")) + # dark blue


  theme_minimal()+
  theme(axis.text.x = element_text(angle = 90, hjust = 1))

```



```
## ## Chunk: Calculate average price per unit by segment
```

```
avg_price_data <- data[, .( Total_Sales = sum(TOT_SALES),
                           Total_Units = sum(PROD_QTY)),
                        by = .(LIFESTAGE, PREMIUM_CUSTOMER)]
# Compute average price per unit
avg_price_data[, Avg_Price := Total_Sales / Total_Units]
# Plot the results
ggplot(avg_price_data, aes(x = LIFESTAGE, y = Avg_Price, fill = PREMIUM_CUSTOMER))+
  geom_bar(stat = "identity", position = "dodge")+
  labs(title = "Average Price per Unit by Lifestage and Premium Status",
       x = "Lifestage", y = "Average Price",
       fill = "Premium Customer")+
  theme_minimal()+
  theme(axis.text.x = element_text(angle = 90, hjust = 1))
```



*# Mainstream midage and young singles and couples are more willing to pay more
per packet of chips compared to their budget and premium counterparts.
This may be due to premium shoppers being more likely to buy healthy snacks
and when they buy chips, this is mainly for entertainment purposes rather than
their own consumption. This is also supported by there being fewer premium
midage and young singles and couples buying chips compared to their mainstream counterparts.*

Chunk: T-Tests

```
# Perform an independent t-test between mainstream vs premium and budget
# midage and young singles and couples
# T-test for Budget vs. Mainstream midage and young singles and couples
t.test(TOT_SALES / PROD_QTY ~ PREMIUM_CUSTOMER,
       data = data,
       subset = LIFESTAGE %in% c("MIDAGE SINGLES/COUPLES", "YOUNG SINGLES/COUPLES")
       & PREMIUM_CUSTOMER %in% c("Budget", "Mainstream"))
```

```
##
## Welch Two Sample t-test
##
## data: TOT_SALES/PROD_QTY by PREMIUM_CUSTOMER
## t = -31.671, df = 23526, p-value <
## 0.00000000000000022
## alternative hypothesis: true difference in means between group Budget and group Mainstream is not eq
## 95 percent confidence interval:
## -0.3738041 -0.3302324
## sample estimates:
## mean in group Budget mean in group Mainstream
```



```
##                                3.687768                4.039786
```

```
# T-test for Budget vs. Premium midage and young singles and couples
t.test(TOT_SALES / PROD_QTY ~ PREMIUM_CUSTOMER,
       data = data,
       subset = LIFESTAGE %in% c("MIDAGE SINGLES/COUPLES", "YOUNG SINGLES/COUPLES")
       & PREMIUM_CUSTOMER %in% c("Budget", "Premium"))
```

```
##
## Welch Two Sample t-test
##
## data: TOT_SALES/PROD_QTY by PREMIUM_CUSTOMER
## t = -2.7694, df = 26724, p-value = 0.005619
## alternative hypothesis: true difference in means between group Budget and group Premium is not equal
## 95 percent confidence interval:
## -0.06347549 -0.01086297
## sample estimates:
## mean in group Budget mean in group Premium
##                3.687768                3.724937
```

```
# The t-test results in a p-value of < 0.00000000000000022, i.e. the unit price
# for mainstream, young and mid-age singles and couples ARE significantly higher
# than that of budget or premium, young and midage singles and couples.
```

```
## ## Chunk: # Find out if the Mainstream, Young singles/couples
## customer segment tend to buy a particular brand of chips
```

```
brand_segment <- data %>%
  group_by(LIFESTAGE, PREMIUM_CUSTOMER, Brand) %>%
  summarise(SegmentBrandQty = sum(PROD_QTY), .groups = "drop")
# Total quantity bought by each segment
segment_total <- data %>%
  group_by(LIFESTAGE, PREMIUM_CUSTOMER) %>%
  summarise(SegmentTotalQty = sum(PROD_QTY), .groups = "drop")
# Total quantity bought per brand
brand_total <- data %>%
  group_by(Brand) %>%
  summarise(BrandTotalQty = sum(PROD_QTY), .groups = "drop")
# Total quantity overall
total_qty <- sum(data$PROD_QTY)
# Merge data
affinity_data <- brand_segment %>%
  left_join(segment_total, by = c("LIFESTAGE", "PREMIUM_CUSTOMER")) %>%
  left_join(brand_total, by = "Brand") %>%
  mutate(
    SegmentBrandShare = SegmentBrandQty / SegmentTotalQty,
    BrandShareOverall = BrandTotalQty / total_qty,
    AffinityScore = SegmentBrandShare / BrandShareOverall
  )
# Focus on Mainstream Young Singles/Couples
affinity_data %>%
  filter(LIFESTAGE == "YOUNG SINGLES/COUPLES", PREMIUM_CUSTOMER == "Mainstream") %>%
```

```
arrange(desc(AffinityScore)) %>%
print (n=21)
```

```
## # A tibble: 20 x 9
##   LIFESTAGE   PREMI~1 Brand Segme~2 Segme~3 Brand~4 Segme~5
##   <chr>      <chr>   <chr>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 YOUNG SING~ Mainst~ Tyrr~    1143    36225    12298 0.0316
## 2 YOUNG SING~ Mainst~ Twis~    1673    36225    18118 0.0462
## 3 YOUNG SING~ Mainst~ Dori~    4447    36225    48331 0.123
## 4 YOUNG SING~ Mainst~ Kett~    7172    36225    79051 0.198
## 5 YOUNG SING~ Mainst~ Tost~    1645    36225    18134 0.0454
## 6 YOUNG SING~ Mainst~ Prin~    4326    36225    48019 0.119
## 7 YOUNG SING~ Mainst~ Cobs     1617    36225    18571 0.0446
## 8 YOUNG SING~ Mainst~ Infu~    2343    36225    27119 0.0647
## 9 YOUNG SING~ Mainst~ Thins    2187    36225    26929 0.0604
## 10 YOUNG SING~ Mainst~ GrnW~    1185    36225    14726 0.0327
## 11 YOUNG SING~ Mainst~ Chee~     651    36225     8747 0.0180
## 12 YOUNG SING~ Mainst~ Smit~    3491    36225    57582 0.0964
## 13 YOUNG SING~ Mainst~ Fren~     143    36225     2643 0.00395
## 14 YOUNG SING~ Mainst~ Chee~     291    36225     5530 0.00803
## 15 YOUNG SING~ Mainst~ RRD      1587    36225    30891 0.0438
## 16 YOUNG SING~ Mainst~ Natu~     710    36225    14106 0.0196
## 17 YOUNG SING~ Mainst~ CCs       405    36225     8609 0.0112
## 18 YOUNG SING~ Mainst~ Sunb~     230    36225     5692 0.00635
## 19 YOUNG SING~ Mainst~ Wool~     873    36225    22333 0.0241
## 20 YOUNG SING~ Mainst~ Burg~     106    36225     2970 0.00293
## # ... with 2 more variables: BrandShareOverall <dbl>,
## #   AffinityScore <dbl>, and abbreviated variable names
## #   1: PREMIUM_CUSTOMER, 2: SegmentBrandQty,
## #   3: SegmentTotalQty, 4: BrandTotalQty,
## #   5: SegmentBrandShare
```

```
# Mainstream, Young singles/couples customer segment prefer to buy chips from
# brands like Tyrrells, Twisties, Kettle, Tostitos and Old.
# Mainstream young singles/couples are 21% more likely to buy Tyrrells chips than
# the other customers.
# Mainstream young singles/couples are 54% less likely to purchase Burger Rings
# compared to the overall customer base.
```

```
## ## Chunk: Find out if our target segment tends to buy larger packs of chips
```

```
# Calculate average PACK_SIZE for Mainstream Young Singles/Couples
# Define your target and all other segment
target_segment <- data %>%
  filter(LIFESTAGE == "YOUNG SINGLES/COUPLES", PREMIUM_CUSTOMER == "Mainstream")
rest_segment <- data %>%
  filter(!(LIFESTAGE == "YOUNG SINGLES/COUPLES" & PREMIUM_CUSTOMER == "Mainstream"))
# Calculate proportion of each pack size in target segment
target_pack_share <- target_segment %>%
  group_by(PACK_SIZE) %>%
  summarise(TargetQty = sum(PROD_QTY)) %>%
  mutate(TotalTarget = sum(TargetQty),
```

```

    TargetProp = TargetQty / TotalTarget)
# Calculate proportion of each pack size in other segments
rest_pack_share <- rest_segment %>%
  group_by(PACK_SIZE) %>%
  summarise(RestQty = sum(PROD_QTY)) %>%
  mutate(TotalRest = sum(RestQty),
         RestProp = RestQty / TotalRest)
# Merge both and calculate affinity score
pack_affinity <- left_join(target_pack_share, rest_pack_share, by = "PACK_SIZE") %>%
  mutate(Affinity_Score = TargetProp / RestProp) %>%
  arrange(desc(Affinity_Score))
# View results
print(pack_affinity)

```

```

## # A tibble: 20 x 8
##   PACK_SIZE Targe~1 Total~2 Targe~3 RestQty Total~4 RestP~5
##   <dbl>     <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1      270     1153   36225 0.0318   10896  434174 0.0251
## 2      380     1165   36225 0.0322   11108  434174 0.0256
## 3      330     2220   36225 0.0613   21779  434174 0.0502
## 4      134     4326   36225 0.119    43693  434174 0.101
## 5      110     3850   36225 0.106    38985  434174 0.0898
## 6      210     1055   36225 0.0291   10907  434174 0.0251
## 7      135      535   36225 0.0148    5677  434174 0.0131
## 8      250      520   36225 0.0144    5549  434174 0.0128
## 9      170     2926   36225 0.0808   35162  434174 0.0810
## 10     150     5709   36225 0.158    70953  434174 0.163
## 11     175     9237   36225 0.255   117230  434174 0.270
## 12     165     2016   36225 0.0557   27035  434174 0.0623
## 13     190      271   36225 0.00748   5402  434174 0.0124
## 14     180      130   36225 0.00359   2634  434174 0.00607
## 15     160      232   36225 0.00640   5372  434174 0.0124
## 16      90      230   36225 0.00635   5462  434174 0.0126
## 17     125      109   36225 0.00301   2621  434174 0.00604
## 18     200      325   36225 0.00897   8100  434174 0.0187
## 19      70      110   36225 0.00304   2745  434174 0.00632
## 20     220      106   36225 0.00293   2864  434174 0.00660
## # ... with 1 more variable: Affinity_Score <dbl>, and
## # abbreviated variable names 1: TargetQty,
## # 2: TotalTarget, 3: TargetProp, 4: TotalRest,
## # 5: RestProp

```

It looks like Mainstream young singles/couples are 27% more likely to purchase a 270g pack of chips compared to the rest of the population but let's dive into what brands sell this pack size.

Check which products are the closest to selling 270g chip packets for targetted marketing
 data[PACK_SIZE == 270, unique(PROD_NAME)]

```
## [1] "Twisties Cheese" "Twisties"
```

```
## ## Conclusion
```

Sales have mainly been due to Budget-older families, Mainstream-young singles/couples,
and Mainstream- retirees shoppers. We found that the high spend in chips for mainstream
young singles/couples and retirees is due to there being more of them than other buyers.
Mainstream, midage and young singles and couples are also more likely to pay more
per packet of chips. This is indicative of impulse buying behaviour. We've also found
that Mainstream young singles and couples are 23% more likely to purchase Tyrrells chips
compared to the rest of the population. The Category Manager may want to increase the
category's performance by off-locating some Tyrrells and smaller packs of chips in
discretionary space near segments where young singles and couples frequent more often to
increase visibilty and impulse behaviour.